

# Big Data Analytics with Hadoop & Apache Spark

## Contents

<b>1</b>	<b>Introduction to Big Data</b>	<b>2</b>
<b>2</b>	<b>Hadoop Ecosystem</b>	<b>2</b>
2.1	What is Hadoop? . . . . .	2
2.2	HDFS Architecture . . . . .	2
2.3	MapReduce Model . . . . .	3
<b>3</b>	<b>Apache Spark</b>	<b>3</b>
3.1	Why Spark is Faster . . . . .	3
3.2	Spark Programming Model . . . . .	3
<b>4</b>	<b>PySpark Examples (Python)</b>	<b>3</b>
4.1	Creating a Spark Session . . . . .	3
4.2	Reading CSV Data . . . . .	4
4.3	Filtering and Selecting Data . . . . .	4
4.4	Group By Aggregation . . . . .	4
4.5	Word Count with RDDs . . . . .	4
<b>5</b>	<b>Spark SQL Example</b>	<b>4</b>
<b>6</b>	<b>Summary</b>	<b>5</b>

## Introduction to Big Data

**Big Data** refers to datasets that are too large, too fast, or too complex for traditional data-processing tools.

The famous **5 V's**:

- **Volume** – massive amounts of data (GB → TB → PB).
- **Velocity** – fast incoming data (real-time streams).
- **Variety** – structured, semi-structured, unstructured.
- **Veracity** – noise, uncertainty in data.
- **Value** – extracting useful insights.

Big Data exists when *your current tools are insufficient* to store or process the data in a reasonable time.

## Hadoop Ecosystem

### What is Hadoop?

Hadoop is an open-source ecosystem enabling distributed storage and processing of massive datasets.

Main components:

- **HDFS** – distributed storage
- **MapReduce** – batch computation model
- **YARN** – cluster resource manager

### HDFS Architecture

HDFS stores files by splitting them into large blocks (default: 128MB) distributed across multiple machines.

Key parts:

- **NameNode** – stores metadata (filesystem tree).
- **DataNode** – stores actual data blocks.
- **Replication** – default 3 copies per block.

## MapReduce Model

MapReduce is a two-step programming model: **Map → Shuffle → Reduce**

Example: Word Count

- Map: Output (word, 1)
- Reduce: Sum counts for each word

## Apache Spark

### Why Spark is Faster

Spark keeps data **in-memory** and executes workflows using an optimized **DAG** engine — often up to 100× faster than classic MapReduce.

Advantages:

- In-memory computing
- Lazy evaluation
- Rich APIs (DataFrames, SQL, MLlib, GraphX)
- Works with HDFS, S3, Kafka

## Spark Programming Model

Types of operations:

- **Transformations** (lazy): `map`, `filter`, `groupByKey`
- **Actions**: `count()`, `collect()`, `show()`

## PySpark Examples (Python)

### Creating a Spark Session

```
from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .appName("BigDataNotes") \
    .master("local[*]") \
    .getOrCreate()
```

## Reading CSV Data

```
df = spark.read.csv("data/people.csv",
                    header=True, inferSchema=True)

df.show()
df.printSchema()
```

## Filtering and Selecting Data

```
adults = df.filter(df.age > 30).select("name", "age")
adults.show()
```

## Group By Aggregation

```
from pyspark.sql import functions as F

city_counts = df.groupBy("city").agg(
    F.count("*").alias("population")
)

city_counts.show()
```

## Word Count with RDDs

```
text = spark.sparkContext.textFile("data/book.txt")

counts = (text.flatMap(lambda line: line.split())
          .map(lambda w: (w.lower(), 1))
          .reduceByKey(lambda a, b: a + b))

counts.take(10)
```

## Spark SQL Example

```
df.createOrReplaceTempView("people")

result = spark.sql("""
    SELECT city, AVG(age) as avg_age
    FROM people
    GROUP BY city
""")

result.show()
```

## Summary

**Hadoop** = distributed storage + batch processing **Spark** = fast, in-memory distributed computation Together they form the backbone of modern Big Data systems.

## Additional Pedagogical Notes and Code Examples

This section contains extra explanations and many code examples in PySpark to help you learn by doing.

### Big Data and HDFS in Practice

Idea: Instead of one “super” machine, we use many cheap machines that work together and store data in HDFS.

Conceptual steps when storing a large file in HDFS:

1. File is split into blocks (e.g. 128MB).
2. Each block is replicated (default 3 copies).
3. Blocks are distributed across several DataNodes.
4. NameNode keeps metadata: which file has which blocks and on which nodes.

Typical HDFS shell commands:

```
# List root directory in HDFS
hdfs dfs -ls /

# Make a directory in HDFS
hdfs dfs -mkdir -p /user/gianna/data

# Upload a local file into HDFS
hdfs dfs -put people.csv /user/gianna/data/

# Show the file content from HDFS
hdfs dfs -cat /user/gianna/data/people.csv

# Remove a file
hdfs dfs -rm /user/gianna/data/people.csv
```

### Working with RDDs (Low-Level API)

RDD = Resilient Distributed Dataset. It is a distributed list of objects that you can transform with `map`, `filter`, `reduceByKey`, etc.

```
# Create RDD from a Python list
numbers = spark.sparkContext.parallelize([1, 2, 3, 4, 5])

# Map: multiply each number by 10
times_ten = numbers.map(lambda x: x * 10)

# Filter: keep only even numbers
even_numbers = times_ten.filter(lambda x: x % 2 == 0)

print("Even numbers:", even_numbers.collect())
```

Word count with RDDs (similar to MapReduce):

```
text_rdd = spark.sparkContext.textFile("data/book.txt")

wc = (text_rdd
      .flatMap(lambda line: line.split())
      .map(lambda w: (w.lower(), 1))
      .reduceByKey(lambda a, b: a + b))

# Show first 20 (word, count) pairs
for word, count in wc.take(20):
    print(word, count)
```

## Creating DataFrames in Different Ways

```
from pyspark.sql import Row

# 1) From a list of Python dicts
data = [
    {"name": "Alice", "age": 30, "city": "Brussels"}, 
    {"name": "Bob", "age": 25, "city": "Paris"}, 
    {"name": "Chloe", "age": 35, "city": "Brussels"}, 
]

df1 = spark.createDataFrame(data)
df1.show()

# 2) From a list of Rows
rows = [
    Row(name="David", age=40, city="Madrid"),
    Row(name="Eva", age=28, city="Athens")
]
df2 = spark.createDataFrame(rows)
df2.show()

# 3) From CSV file with inferred schema
df3 = (spark.read
        .option("header", True)
        .option("inferSchema", True)
        .csv("data/people.csv"))

df3.show()
df3.printSchema()
```

## DataFrame: Select, Filter, New Columns

```
from pyspark.sql import functions as F

df = df3 # assuming df3 from previous example

# Select a subset of columns
df.select("name", "age").show()

# Filter rows: age > 30
df.filter(df.age > 30).show()

# Multiple conditions (AND / OR)
df.filter((df.age > 25) & (df.city == "Brussels")).show()

# Add a new column (computed)
df2 = df.withColumn("age_in_10_years", df.age + 10)
df2.show()

# Rename a column
df_renamed = df2.withColumnRenamed("age", "current_age")
df_renamed.show()
```

## GroupBy & Aggregations (More Examples)

```
from pyspark.sql import functions as F

# Count how many people per city
df.groupBy("city").agg(
    F.count("*").alias("n_people")
).show()

# Average and max age per city
df.groupBy("city").agg(
    F.avg("age").alias("avg_age"),
    F.max("age").alias("max_age")
).show()

# Filter on aggregated result (e.g. cities with avg_age > 30)
agg = df.groupBy("city").agg(
    F.avg("age").alias("avg_age")
)

agg.filter(agg.avg_age > 30).show()
```

## Joins Between DataFrames

Imagine two datasets:

- `customers(id, name, city)`
- `orders(order_id, customer_id, amount)`

```
customers = spark.read.csv("data/customers.csv",
                           header=True, inferSchema=True)
orders = spark.read.csv("data/orders.csv",
                       header=True, inferSchema=True)

# Inner join on customer id
joined = orders.join(customers,
                     orders.customer_id == customers.id,
                     "inner")

joined.select("order_id", "name", "amount", "city").show()

# Left join: keep all orders even if no customer match
left_join = orders.join(customers,
                        orders.customer_id == customers.id,
                        "left")

left_join.show()
```

## User-Defined Functions (UDFs)

Use a UDF when a transformation cannot easily be expressed with built-in functions.

```

from pyspark.sql.functions import udf, col
from pyspark.sql.types import StringType

# Python function
def categorize_age(age):
    if age is None:
        return "unknown"
    if age < 18:
        return "child"
    elif age < 30:
        return "young adult"
    elif age < 60:
        return "adult"
    else:
        return "senior"

# Convert Python function to UDF
age_category_udf = udf(categorize_age, StringType())

df_with_cat = df.withColumn("age_category",
                             age_category_udf(col("age")))

df_with_cat.select("name", "age", "age_category").show()

```

## Window Functions (Running Totals, Rankings)

Window functions let you compute things like “running sum”, “rank per group”, “previous value”, etc.

Example: sales per day and running total per city.

```

from pyspark.sql import functions as F
from pyspark.sql.window import Window

sales = spark.read.csv("data/daily_sales.csv",
                      header=True, inferSchema=True)
# Columns: date, city, amount

w = Window.partitionBy("city").orderBy("date")

sales_with_running = sales.withColumn(
    "running_total",
    F.sum("amount").over(w)
)

sales_with_running.show()

```

## Spark SQL: More Complex Queries

```
df.createOrReplaceTempView("people")

# Cities with at least 2 people and average age > 30
result = spark.sql("""
    SELECT city,
           COUNT(*) AS n_people,
           AVG(age) AS avg_age
      FROM people
     GROUP BY city
    HAVING COUNT(*) >= 2 AND AVG(age) > 30
   ORDER BY avg_age DESC
""")  
result.show()
```

## Machine Learning with MLlib: Classification

Workflow idea: **DataFrame** → **features vector** → **model fit** → **predictions**.

Example: simple logistic regression.

```
from pyspark.ml.feature import VectorAssembler, StringIndexer
from pyspark.ml.classification import LogisticRegression
from pyspark.ml import Pipeline

# Suppose df has: label (0/1), age, income, balance
data = spark.read.csv("data/bank.csv",
                      header=True, inferSchema=True)

# Features into one vector column
assembler = VectorAssembler(
    inputCols=["age", "income", "balance"],
    outputCol="features"
)

# Model
lr = LogisticRegression(featuresCol="features",
                        labelCol="label")

pipeline = Pipeline(stages=[assembler, lr])

train, test = data.randomSplit([0.7, 0.3], seed=42)

model = pipeline.fit(train)

predictions = model.transform(test)

predictions.select("age", "income", "balance",
                   "label", "prediction", "probability").show(20,
                                                 truncate=False)
```

## MLlib: Regression Example

```
from pyspark.ml.regression import LinearRegression

housing = spark.read.csv("data/housing.csv",
                        header=True, inferSchema=True)
# Example columns: price, rooms, size, distance_to_center

assembler = VectorAssembler(
    inputCols=["rooms", "size", "distance_to_center"],
    outputCol="features"
)

housing_vec = assembler.transform(housing)

train, test = housing_vec.randomSplit([0.8, 0.2], seed=1)

lr = LinearRegression(featuresCol="features",
                      labelCol="price")

lr_model = lr.fit(train)

print("Coefficients:", lr_model.coefficients)
print("Intercept:", lr_model.intercept)

pred = lr_model.transform(test)
pred.select("rooms", "size", "distance_to_center",
            "price", "prediction").show(10)
```

## MLlib: K-Means Clustering

```
from pyspark.ml.clustering import KMeans

points = spark.read.csv("data/points2d.csv",
                       header=True, inferSchema=True)
# Columns: x, y

assembler = VectorAssembler(
    inputCols=["x", "y"],
    outputCol="features"
)

points_vec = assembler.transform(points)

kmeans = KMeans(k=3, seed=1)
model = kmeans.fit(points_vec)

centers = model.clusterCenters()
print("Cluster centers:")
for c in centers:
    print(c)

clustered = model.transform(points_vec)
clustered.select("x", "y", "prediction").show(20)
```

## Spark Structured Streaming: Simple Example

Structured Streaming treats streaming data as an unbounded table. We write queries that are continuously updated.

Example: word count from a socket.

```

from pyspark.sql import functions as F

# In a terminal, run: nc -lk 9999
# Then type lines of text.

lines = (spark.readStream
         .format("socket")
         .option("host", "localhost")
         .option("port", 9999)
         .load())

words = lines.select(F.explode(F.split(lines.value, " ")).alias("word"))

word_counts = words.groupBy("word").count()

query = (word_counts.writeStream
          .outputMode("complete")
          .format("console")
          .start())

query.awaitTermination()

```

## Mini ETL Pipeline Example

ETL = Extract → Transform → Load. Spark is often used to build ETL pipelines on Big Data.

Goal: Read raw CSV of transactions, clean data, aggregate by customer, and write result as Parquet.

```

from pyspark.sql import functions as F

# 1. Extract
raw = spark.read.csv("data/transactions.csv",
                      header=True, inferSchema=True)
# Columns: customer_id, date, amount, country

# 2. Transform
clean = (raw
         .filter(raw.amount.isNotNull())
         .filter(raw.customer_id.isNotNull()))

aggregated = (clean
               .groupBy("customer_id")
               .agg(
                   F.count("*").alias("n_tx"),
                   F.sum("amount").alias("total_spent"),
                   F.avg("amount").alias("avg_ticket")
               ))
               .show()

# 3. Load (save)
(aggregated.write
     .mode("overwrite")
     .parquet("output/customers_aggregated"))

# You can later read:
result = spark.read.parquet("output/customers_aggregated")
result.show()

```

Try to re-run the ETL pipeline with different filters or new columns. You learn Spark best by experimenting and breaking things yourself.